

Concurrency and Parallelism 2022-23

Lab 2: Calculation of π Using the Monte Carlo Method using C

Hervé Paulino & Alex Davidson

March 2023

Abstract

With this second lab assignment, you will learn/recall some basic concepts about developing concurrent/parallel programs using the C programming language and the `pThreads` interface and library.

1 Recall: Monte Carlo Methods

Monte Carlo methods [1] can be thought of as statistical simulation methods that utilize a sequence of random numbers to perform some kind of simulation. The name Monte Carlo was coined by Nicholas Constantine Metropolis (1915-1999) and inspired by Stanislaw Ulam (1909-1986), because of the similarity of statistical simulation to games of chance, and because Monte Carlo is a centre for gambling and games of chance. In a typical two-dimensional process, one computes the number of points within a shape A that lies inside a larger shape R . The ratio of the number of points that fall inside A to the total number of points tried is equal to the ratio of the two areas (or volume in a three-dimensional process). The accuracy of the ratio ρ depends on the number of points used, with more points leading to a more accurate value.

1.1 Estimating the value of π

Figure 1 shows a circle with radius $r = 1$ inscribed within a square. The area of the circle is

$$A_{\bigcirc} = \pi \cdot r^2 = \pi \cdot 1^2 = \pi$$

and the area of the square is

$$A_{\square} = (2 \cdot r)^2 = 2^2 = 4.$$

The ratio ρ of the area of the circle to the area of the square is

$$\rho = \frac{A_{\bigcirc}}{A_{\square}} = \frac{\pi}{4} = 0.7853981634.$$

A simple Monte Carlo simulation to approximate the value of π can first estimate ρ and then π . Given ρ then

$$\pi = \rho \times 4 = 0.7853981634 \times 4 = 3.1415926536.$$




Figure 1: A circle within a square.

One particularly simple way to compute ρ is to pick random points in the square and count how many of them lie inside the circle (Figure 2).

A Monte Carlo simulation to approximate the value of π will then involve randomly selecting points $(x_i, y_i)_{i=1}^n$ in the unit square and determining the ratio $\rho = \frac{m}{n}$, where n is the total number of points and m is number of points that satisfy $x_i^2 + y_i^2 \leq 1$ (see Figure 3).

In a simulation of sample size $n = 1000$, there were 787 points satisfying that equation. Using this data we obtain

$$\rho = \frac{m}{n} = \frac{787}{1000} = 0.787$$

and

$$\pi = A_{\bigcirc} = \rho \times A_{\square} = \rho \times 4 = 0.787 \times 4 = 3.148.$$

If the numbers used in the Monte Carlo simulation are random, every time a simulation is made using the same sample size n , it will arrive at a slightly different value. The values converge very slowly with order $O(n^{-1/2})$. This property is a consequence of the Central Limit Theorem [2].

You may find a web simulation of this method at <https://academo.org/demos/estimating-pi-monte-carlo/>.




Figure 2: Ratio between areas of circle and square.




Figure 3: Monte Carlo method.

2 Lab Work

Open the following link, and log on to GitHub if requested

<https://classroom.github.com/a/Z2b6gvy0>

and then follow the instructions below.

In the give C program, you are expected to implement the functions

```
void *worker (void *arg)
```

and

```
double calculate_pi (long nPoints, int nThreads)
```

You may compile the C source file with:

```
gcc -o approxPi approxPi.c -lpthread
```

2.1 Parallel Version of the Monte Carlo Simulation

Design and implement a C (parallel) program named `approxPi` that makes use of the `pThreads` library to approximate the value of π by using the Monte Carlo method. In the command line, the program accepts 2 options and must receive 2 mandatory arguments.

The usage is:

```
./approxPi [-d] [-s] nPoints nThreads
    -d: Enable debugging more info is printed on the screen.
    -s: Silent mode just print numbers with no explanation
        of their meaning.
    nPoints: The total number of points to be used in the simulation
        (to be divided by the working threads).
    nThreads: The number of working threads.
```

Examples below (the \$ is the shell prompt and is not part of the command):

```
$ ./approxPi 100000000 8
Pi: 3.312500000
Real Time: 0.261323
CPU Time: 0.261323
```

```
$ ./approxPi -s -d 100000000 8
nPoints: 100000000
nThreads: 8
Worker: 12500000
Worker: 12500000
Worker: 12500000
Worker: 12500000
Worker: 12500000
Worker: 12500000
Worker: 12500000
Worker: 12500000
Win: 12500000
Win: 12500000
Win: 10156250
Win: 10937500
Win: 10937500
Win: 11718750
Win: 10937500
Win: 11718750

nPnt      nThr  Real      CPU      Pi
100000000  8      0.267139  1.902000  3.656250000
```

```
$ ./approxPi -s 100000000 8
100000000 8 0.250770 1.889024 3.437500000
```

You may use the script `run_tests.sh` given in the repository to help executing the tests. The options given to the script will be passed to the `approxPi` program. For example, you may run it as:

```
$ ./run_tests.sh -s
```

Remark: Remember to keep using these simple lab exercises to learn how to use *git*. Use the given repository (and the give starting code) to manage the versioning of your code. Remember to do frequent commits with clear commit messages.

2.2 Analysing Results

Once your program is running correctly:

1. Try your program with some values (e.g., 100, 1000, 10000, etc) and build a plot with the execution time (t) as a function of the number of points (n) and of the number of processors/threads (p). **Remember to run each configuration in a way that allows you to establish scientific certainty of your results.**¹
2. Answer the following questions:
 - The parallel version with two threads is faster than with one?
 - When you double the number of threads does it take half the time? More? Less?
 - Which of your implementations is faster? C or Java? Why?

2.3 Extension: Performance Profiling

To go deeper into analysing why the performance of your program is a certain way, you will need to spend some time profiling your compiled code to see which functions take the longest. You can a variety of tools for doing, for example the CLion IDE allows you to profile certain things directly. On the command line, there are a variety of tools that can help you profile and visualise your program at a functional level (such as `perf`, `dtrace`, etc.) This blog post explains some different ways of doing this.

Final Note

There is a large bibliography of work on this problem, the corresponding Monte Carlo simulation, and its implementation in a variety of programming languages. Please understand that the objective of this homework is for you to learn about concurrency and this lab assignment is **not** graded. So, there is no value in cheating by searching the web for working solutions of this problem! *Just be honest with yourself and develop your own solution, or just ignore it!*

Acknowledgments

The text from the first two sections is an adaptation of the text found at <http://mathfaculty.fullerton.edu/mathews/n2003/montecarlopimod.html>.

References

- [1] A. M. Johansen. “Monte Carlo Methods”. In: *International Encyclopedia of Education (Third Edition)*. Ed. by Penelope Peterson, Eva Baker, and Barry McGaw. Third Edition. Oxford: Elsevier, 2010, pp. 296–303. ISBN: 978-0-08-044894-7. DOI: 10.1016/B978-0-08-044894-7.01543-8.
- [2] S. L. Zabell. “Alan Turing and the Central Limit Theorem”. In: *The American Mathematical Monthly* 102.6 (1995), pp. 483–494. ISSN: 00029890, 19300972. URL: <http://www.jstor.org/stable/2974762>.

¹For example, think about the number of times that you run it, and under what conditions the program should run.